

ControlPlane.ai — Business Proposal

Enterprise AI Governance, Responsible Agent Security & Safety Gateway

1. Executive Summary

As enterprises rapidly adopt Large Language Models (LLMs) and Autonomous AI Agents, governance, compliance, financial risk control, and safety have emerged as mandatory operational imperatives. **ControlPlane.ai** is an enterprise-grade AI governance and safety gateway that sits directly between applications and upstream AI models or tool-calling agents. Rather than relying on unmonitored model outputs, ControlPlane programmatically inspects inputs/outputs, enforces custom tenant policies, redacts sensitive PII, verifies factual claims against knowledge repositories, tracks multi-turn data taint, and gates high-consequence tool calls before real-world execution.

2. The Market Pain & Core Problem

Enterprise deployments of Generative AI face three critical vulnerability vectors:

- **Hallucination & Financial Liability:** Unverified claims made by customer service bots (e.g. promising non-existent refunds or incorrect pricing) can lead to direct legal liability.
- **Data Contamination & Tool Exploitation:** Tainted or unverified claims generated early in a conversation can carry over into downstream tool calls (e.g., automated execution of ``issue_refund``), executing real monetary transactions.
- **Compliance & Regulatory Risk:** Tightening regulations like the EU AI Act and NIST AI RMF mandate cryptographically verifiable audit trails, PII redaction, and human-in-the-loop oversight.

3. ControlPlane.ai Solution Architecture

Governance Pillar	ControlPlane Technical Capability	Business Value Impact
Adaptive Scrutiny	Tier 0 cheap heuristic gate + Tier 1 claim verification pipeline.	Reduces API cost by routing non-critical requests through low-latency gates.
Surgical Remediation	Per-claim remediation (Hedge, Redact, Remove, Cite, Escalate).	Preserves positive user experience without dropping full responses unnecessarily.
Multi-Turn Taint Tracking	Conversation state tracking linking unverified claims to downstream tool parameters.	Prevents unauthorized financial transactions and rogue agent actions.
Immutable Audit Ledger	Cryptographically hash-chained transaction logging.	Ensures 100% audit readiness for regulatory compliance.
Dynamic Risk Appetite	Live slider control + human reviewer recalibration feedback loop.	Empowers compliance officers to adjust risk appetite live without code deployments.

4. Market Opportunity & Business Model

The total addressable market (TAM) for Enterprise AI Trust, Risk, and Security Management (AI TRiSM) is projected to exceed **\$35 Billion by 2028**.

Monetization Architecture:

- 1. **SaaS Gateway Tier:** Consumption-based model per 1M governed tokens + active managed tenant policies.
- 2. **Enterprise VPC / On-Prem License:** Flat annual enterprise license per gateway instance with dedicated compliance connectors.
- 3. **Accenture Managed Services Integration:** Partner integration embedding ControlPlane into Accenture's AI Transformation offerings.

5. Financial Projections (USD)

Metric	Year 1	Year 2	Year 3
Enterprise Clients	15	60	220
Annual Recurring Revenue (ARR)	\$1.2M	\$5.8M	\$24.5M
Gross Margin (%)	78%	82%	86%

6. Strategic Alignment with Accenture

ControlPlane provides Accenture with a high-value proprietary governance asset to differentiate its AI Consulting and System Integration practices. By offering pre-built compliance policies for regulated industries (Healthcare, Banking, Public Sector), Accenture can accelerate enterprise AI deployments while mitigating risk.

7. Product Roadmap & Milestones

- **Q3 2026:** Production gateway rollout, Kubernetes Operator release, and automated NIST AI RMF compliance reporting.
- **Q4 2026:** Integration with Enterprise SSO (Okta/Azure AD), Hardware Security Module (HSM) hash signing, and fine-tuned NLI verification models.

8. Project Team & Contributors

- **Mohith Chandra:** Team Lead
- **Monal Gupta:** Core NLP & Fullstack Developer
- **Veda Vikas:** System Architect & Developer